PCA for extremes

Sam Morris[1], Brian J Reich[1], Emeric Thibauld[2], and Dan Cooley[2]

February 16, 2016

Abstract

words...

Key words: Max-stable process.

---

[1]North Carolina State University

[2]Colorado State University

1 Introduction

2 Model

Let Y_{it} be the observation at location $\mathbf{s}_i$ for $i \in \{1, ..., n_s\}$ and time $t \in \{1, ..., n_t\}$. To focus attention on the extreme values, we consider data above a threshold T. The marginal distribution of Y_{it} is then determined by the probability of exceeding the threshold and the distribution of the excursions. Denote the exceedance probability as $\text{Prob}[Y_{it} > T] = p_{it}$. Extreme value theory says that for sufficiently large T the excursion distribution can be approximated using a generalized Pareto distribution (GPD). Therefore we model $Y_{it}|Y_{it} > T \sim \text{GDP}(\sigma_{it}, \xi)$, where the GDP scale and shape parameters are denoted $\sigma_{it} > 0$ and ξ, respectively.

spectral, max-linear...finally we settle on... Spatial extremal dependence is captured using a max-stable copula (define). Let Z_{it} be a max-stable process with Fréchet marginal distributions (define GEV etc...). Our objective is to identify a low-rank model for spatial dependence in Z_{it}. Decompose Z_{it} as $Z_{it} = \theta_{it}\varepsilon_{it}$ where θ_{it} is a spatial process and $\varepsilon_{it} \stackrel{iid}{\sim} \text{GEV}(1, \alpha, \alpha)$ is a nugget. The spatial component is written as a combination of L basis functions B_{il}

$$\theta_{it} = \left(\sum_{l=1}^{L} B_{il}^{1/\alpha} A_{lt} \right)^{\alpha}. \tag{1}$$

If $B_{il} > 0$, $\sum_{l=1}^{L} B_{il} = 1$, and the A_{lt} have positive stable (PS) distribution $A_{lt} \sim \text{PS}(\alpha)$ (define), then Z_{it} is max-stable and has Fréchet marginal distributions.

The Z_{it} are conditionally independent given the spatial random effects, with conditional distribution $Z_{it}|\theta_{it} \sim$. As a result, the likelihood is $Y_{it}|\theta_{it} \stackrel{indep}{\sim} g(y; \theta_{it}, p_{it}, \sigma_{it}, \xi)$ where

$$g(y; \theta, p, \sigma, \xi) = \tag{2}$$

Therefore, the likelihood factors across observations which is computationally convenient. Marginalizing over the random effect θ_{it} induces extremal spatial dependence in the Z_{it}, and thus the Y_{it}.

Spatial dependence can be summarized by the extremal coefficient (EC) $\vartheta_{ij} \in [1, 2]$, where

$$\text{Prob}(Z_{it} < c, Z_{jt} < c) = \text{Prob}(Z_{it} < c)^{\vartheta_{ij}}. \tag{3}$$

For the PS random effects model the EC has the form

$$\vartheta_{ij} = \sum_{l=1}^{L} \left(B_{il}^{1/\alpha} + B_{jl}^{1/\alpha} \right)^{\alpha}. \tag{4}$$

In particular, $\vartheta_{ii} = 2^{\alpha}$ for all i. Since $\sum_{l=1}^{L} B_{il} = 1$ for all i, we have $\sum_{l=1}^{L} (\sum_{i=1}^{n_s} B_{il}/n_s) = 1$. Therefore, the relative contribution of term l can be measured by

$$v_l = \sum_{i=1}^{n_s} B_{il}/n_s, \tag{5}$$

with $\sum_{l=1}^{L} v_l = 1$. The order of the terms is arbitrary, and so we assume without loss of generality that $v_1 \geq ... \geq v_L$.

3   Estimating the extremal coefficient function

In this section we develop an algorithm to estimate the spatial dependence parameter α and the $n_s \times L$ matrix $\mathbf{B} = \{B_{il}\}$. Given these parameters, we plug them into our model and proceed with Bayesian analysis as described in Section 4. Our algorithm has the following steps:

(1) Obtain an initial estimate of the extremal coefficient for each pair of locations, $\hat{\vartheta}_{ij}$.

(2) Spatially smooth these initial estimates $\hat{\vartheta}_{ij}$ using kernel smoothing to obtain $\tilde{\vartheta}_{ij}$.

(3) Estimate the spatial dependence parameters by minimizing the difference between model-based coefficients, ϑ_{ij}, and smoothed coefficients, $\tilde{\vartheta}_{ij}$.

To estimate the spatial dependence we first remove variation in the marginal distribution. Let $U_{it} = \sum_{k=1}^{n_t} I[Y_{ik} < Y_{it}]/n_t$, so that the U_{it} are approximately uniform at each location. Then for

some extreme probability $q \in (0, 1)$, solving (3) suggest the estimate

$$\hat{\vartheta}_{ij}(q) = \frac{\log[Q_{ij}(q)]}{\log(q)}, \tag{6}$$

where $Q_{ij}(q) = \sum_{t=1}^{n_t} I[U_{it} < q, U_{jt} < q]/n_t$ is the sample proportion of the time points at which both sites are less then q. Since all large q give valid estimates, we average over a grid of q with $q_1 < \cdots < q_{n_q}$

$$\hat{\vartheta}_{ij} = \frac{1}{n_q} \sum_{j=1}^{n_q} \hat{\vartheta}_{ij}(q_j). \tag{7}$$

Assuming the true B_{il} are smooth over space, the initial estimates $\hat{\vartheta}_{ij}$ can be improved by smoothing. Let

$$\tilde{\vartheta}_{ij} = \frac{\sum_{u=1}^{n_s} \sum_{v=1}^{n_s} w_{iu} w_{jv} \hat{\vartheta}_{uv}}{\sum_{u=1}^{n_s} \sum_{v=1}^{n_s} w_{iu} w_{jv}}, \tag{8}$$

where $w_{iu} = \exp(-\phi||\mathbf{s}_i - \mathbf{s}'_u||^2)$ is the Gaussian kernel function with bandwidth ϕ. The elements $\hat{\vartheta}_{ii}$ do not contributed any information as $\hat{\vartheta}_{ii} = 1$ for all i by construction. To eliminate the influence of these estimates we set $w_{ii} = 0$. However, this approach does give imputed values $\tilde{\vartheta}_{ii}$, which provides information about small-scale spatial variability.

The dependence parameters are estimated by comparing estimates $\tilde{\vartheta}_{ij}$ with the model-based values ϑ_{ij}. For all i, $\vartheta_{ii} = 2^\alpha$, and therefore we set α to $\hat{\alpha} = \log_2(\sum_{i=1}^{n_s} \tilde{\vartheta}_{ii}/n_s)$. Given $\alpha = \hat{\alpha}$, it remains to estimate $\mathbf{B}$. These estimate $\hat{\mathbf{B}}$ is taken as the minimizer of

$$m(\mathbf{B}) = \sum_{i<j} \left(\tilde{\vartheta}_{ij} - \vartheta_{ij} \right)^2 = \sum_{i<j} \left\{ \tilde{\vartheta}_{ji} - \sum_{l=1}^{L} [B_{il}^{1/\hat{\alpha}} + B_{jl}^{1/\hat{\alpha}}]^{\hat{\alpha}} \right\}^2 \tag{9}$$

under the restrictions that $B_{il} \geq 0$ for all i and l and $\sum_{l=1}^{L} B_{il} = 1$ for all i.

The order of the B_{il} is not defined. Therefore, we sort the terms so that $v_1 > \ldots > v_L$.

4 Implementation details

The model has three tuning parameters: the quantile threshold q, the kernel bandwidth ϕ, and the number of terms L. How to pick? Say $q = 0.95$ or whatever seems to give GPD marginals. ϕ is

3

61 something reasonable. For L, we start small and increase until the smallest proportion v_L is less

62 than, say 0.05.

63   Given the estimates of α and $\mathbf{B}$, the hierarchical model is

$$
\begin{aligned}
Y_{it}|\theta_{ij} &\overset{indep}{\sim} g(y; \theta_{it}, p_{it}, \sigma_{it}, \xi) \qquad\qquad (10)\\
\theta_{it} &= \left(\sum_{l=1}^{L} \hat{B}_{il}^{1/\hat{\alpha}} A_{lt} \right)^{\hat{\alpha}} \quad \text{where} \quad A_{lt} \overset{iid}{\sim} PS(\hat{\alpha})\\
\text{logit}(p_{it}) &= \mathbf{X}_{it}^T \boldsymbol{\beta}_1 \quad \text{and} \quad \log(\sigma_{it}) = \mathbf{X}_{it}^T \boldsymbol{\beta}_2
\end{aligned}
$$

64 where g is given in (2) and $\mathbf{X}_{it} = (X_{it1}, ..., X_{itp})^T$ is a vector of spatiotemporal covariates. To

65 complete the Bayesian model, we select independent normal priors with mean zero and variance

66 100 for the components of $\boldsymbol{\beta}_1$ and $\boldsymbol{\beta}_2$ and standard normal prior for ξ.

67   We estimate parameters $\Theta = \{A_{lt}, \boldsymbol{\beta}_1, \boldsymbol{\beta}_2, \xi\}$ using Markov chain Monte Carlo. Details...

5   Data analysis

69 The dataset used for our application is composed of yearly acreage burned due to wildfires for each

70 county in Georgia from $1965 - 2014$ (`http://weather.gfc.stat.ga.us/FireData/`).

71   **Plots of a couple year's of data**

72   Although some counties contain acres burned for years prior to 1965, we choose to start at 1965

73 because that is the first year for which data are available for all counties. We estimate the extremal

74 coefficient function $\hat{\theta}_{ij}$ by setting $q_1 = 0.90$ and using $n_q = 100$. We set $q_1 = 0.90$ because we

75 only have 49 years of data for each site.

5.1 Results

5.2 Model checking and sensitivity analysis

6 Conclusions

Acknowledgements

References